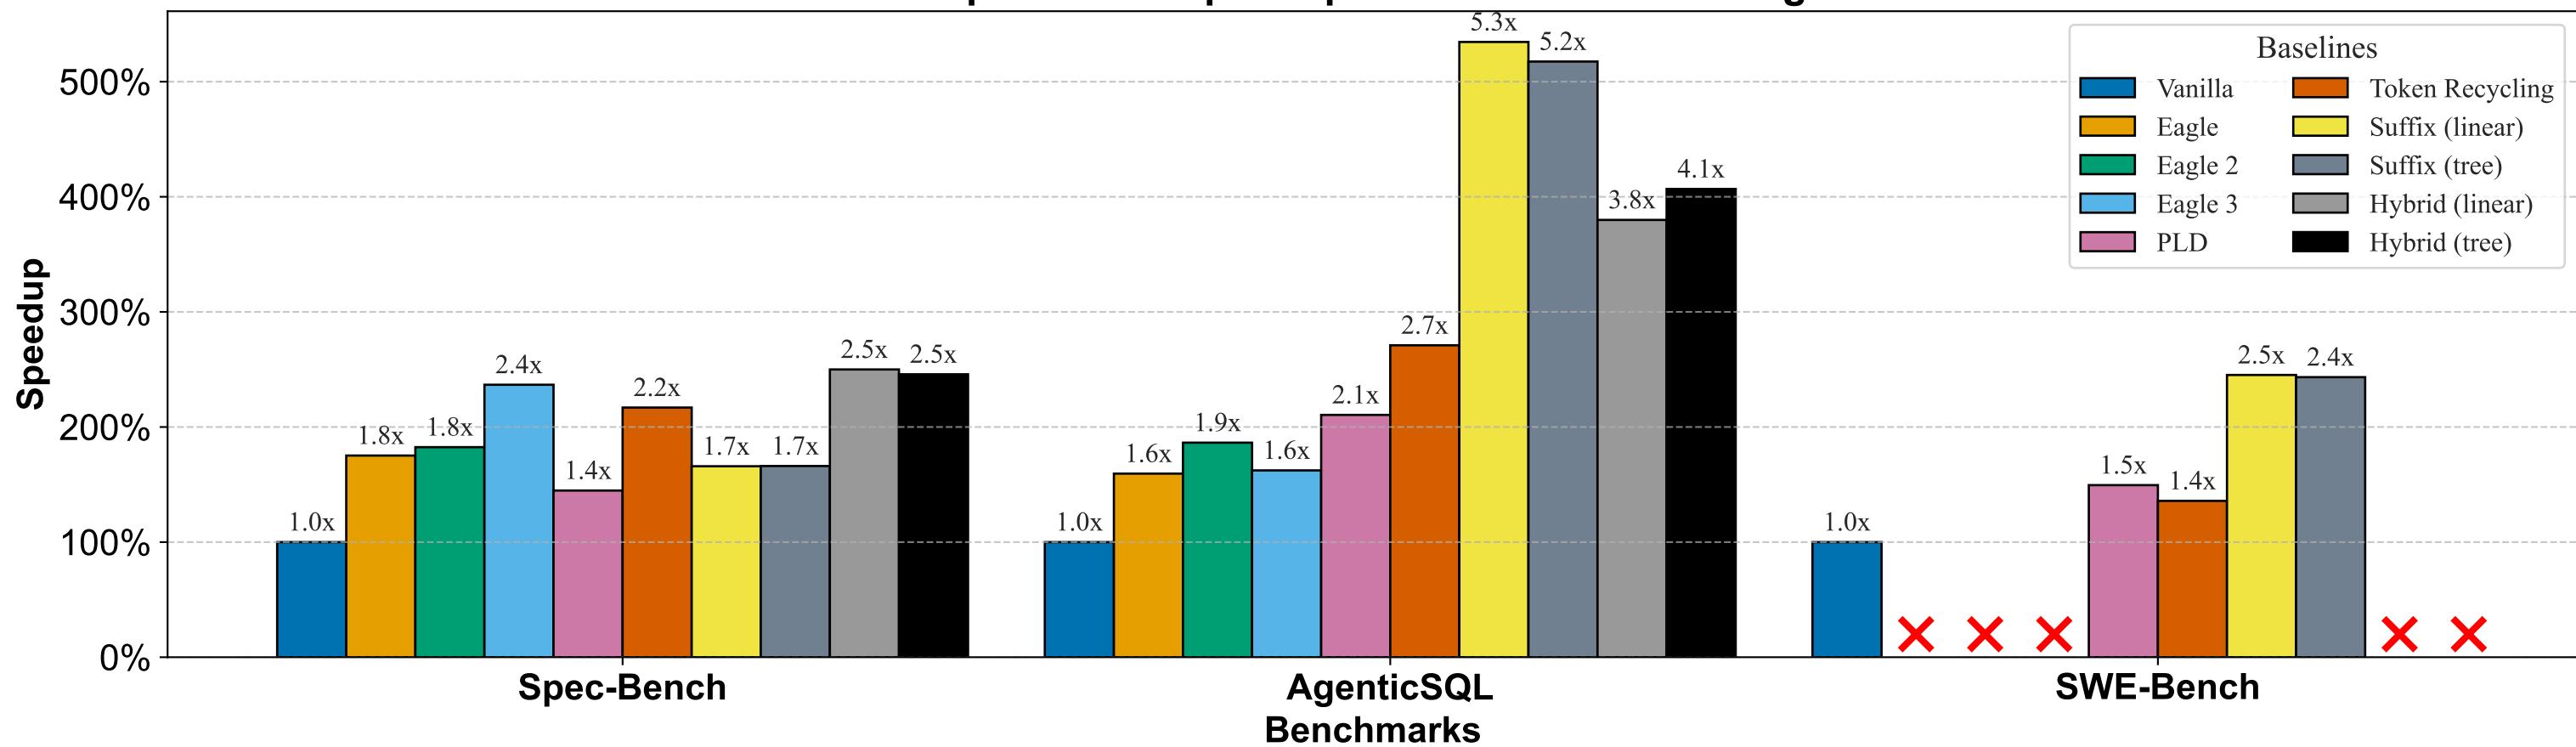


# Speculative Speedups over Vanilla Decoding



## Mean Accepted Tokens per Step

